

DAG scheduling

Guoli Ding and Rod Tohid

September 19, 2019

The inputs.

- A positive integer m , indicating that there are m machines available.
- An acyclic directed graph $D = (V, E)$, known as a DAG. Here a vertex represents a job and a directed edge $e = uv$ indicates that job v cannot start before job u finishes.
- A nonnegative weight function ℓ on $V \cup E$, where ℓ_v indicates the amount of time required to finish job v , and ℓ_{uv} indicates the amount of time required for v to access u if u, v are located on different machines.

Assumptions.

- A machine can process at most one job at a time.
- If a machine starts a job then it has to finish the job before it can start another job. In particular, this implies that a job can only be completed by one machine.
- For each job v , the weight ℓ_v is the same on every machine. That is, a job can be done by any machine with the same amount of time, or simply put, the machines are identical.
- For any edge $e = uv$, the weight ℓ_e remains the same no matter which two machines process u and v . In other words, the communication channels are identical.

The problem.

- Decide which job goes to which machine.
- Decide a starting time for each job.
- Minimize the final completion time of all jobs.

Linear programming formulation. Let M_1, \dots, M_m denote the machines and let v_1, \dots, v_n denote the jobs. We introduce the following variables.

- t_i : the starting time for v_i
- T : final finishing time
- x_{ij} : the assignment indicator, which is 1 if job v_i is assigned to machine M_j , and is 0 if otherwise
- y_{ij} : the time line indicator, which is 1 if $t_i \leq t_j$ and is 0 if $t_i > t_j$
- L : defined as the sum of ℓ over $V \cup E$, which is the total amount of involved time

The following is the formulation.

min T

subject to:

$$x_{i1} + x_{i2} + \dots + x_{im} = 1, \quad \forall v_i \in V \tag{1}$$

$$t_j - t_i - \ell_i \geq \ell_{ij} \max\{x_{ik} - x_{jk} : k = 1, \dots, m\}, \quad \forall v_i, v_j \in E \tag{2}$$

$$t_j - t_i - \ell_i \geq L(\max\{x_{ik} + x_{jk} - 2 : k = 1, \dots, m\} + y_{ij} - 1), \quad \forall v_i, v_j \in V \text{ with } i < j \tag{3a}$$

$$t_i - t_j - \ell_j \geq L(\max\{x_{ik} + x_{jk} - 2 : k = 1, \dots, m\} - y_{ij}), \quad \forall v_i, v_j \in V \text{ with } i < j \tag{3b}$$

$$T \geq \max\{t_i + \ell_i : i = 1, \dots, n\} \tag{4}$$

$$t_i \geq 0; \quad x_{ij} \in \{0, 1\}; \quad y_{ij} \in \{0, 1\}$$

Note that the number of variables is $n + 1 + mn + \binom{n}{2}$ and the number of inequalities is $m(n^2 + |E| - n) + 3n$. Some of the inequalities/variables are redundant and thus can be removed.

Greedy Algorithm.

- Keep track of the set R of ready jobs (with cleared dependencies)
- As soon as a machine becomes available,
 - update R
 - assign a job from R to this machine
- Repeat until all jobs are scheduled.

REMARKS.

- At any point in time, let F denote the set of jobs that are finished (so $F = \emptyset$ in the very beginning). Then R consists of jobs v such that no edge of $D \setminus F$ is directed to v . That is, the in-degree of v is zero in $D \setminus F$.
- When assigning $v \in R$ to an available machine, there are many choices. We could develop different principles for making such a choice and let experiments decide which is better –this is further research.

An application. Suppose we have a user program \mathbb{P} , by which we mean a sequence of matrix operations (an example of \mathbb{P} is illustrated below). Program \mathbb{P} can be expressed by a computation graph \mathbb{G} (which is essentially the execution tree) and this graph is a DAG (as illustrated below). If we want to process \mathbb{P} in m machine then we have exactly the DAG scheduling problem.

